

Hiring Task: AI-Based UPSC Answer Sheet Evaluation with RAG

Objective

Design and build an **AI-powered automated grading workflow** that evaluates handwritten UPSC answers using **vision and language models**, and incorporates a **RAG (Retrieval-Augmented Generation)** system grounded in our official SOP Evaluation document.

Task Overview

Step 1: Input Handling

- Accept input as **PDFs or images of handwritten answer sheets**.
- These documents contain:
 - A visible **question**
 - The **student's answer**, structured across **Introduction, Body, Conclusion**
 - **Marking scheme** (e.g., "15 marks") embedded in the question text

Step 2: Workflow Automation

You must build an **end-to-end pipeline** that includes:

a. OCR Extraction

- Extract handwritten content using a robust OCR system (e.g., **Google Vision API**, **PaddleOCR**, **Amazon Textract**, or vision-LMs like GPT-4o).
- Segment extracted text into structured fields: **question**, **name**, **introduction**, **body**, **conclusion**.

b. RAG-based Evaluation

- Implement a **Retrieval-Augmented Generation system**:
 - The retriever must reference a knowledge base built from the shared PDF: **SOP_Evaluation_Test_Series_Mentorship_Hindi.pdf**
 - You must chunk, embed, and store this SOP in a **vector database** like Pinecone, Weaviate, or Qdrant.
 - When evaluating a student's response, the generator (LLM) should:
 - Retrieve relevant SOP snippets (in Hindi) per section (Intro, Body, Conclusion, etc.)
 - Use them as grounding context while generating feedback and assigning marks.

c. Evaluation Criteria

Evaluate each section based on:

- **Content Adequacy**
- **Structure (IBC format)**
- **Depth of Argument & Analysis**
- **Use of Examples / Data**
- **Presentation**
- **Language and Grammar**
- **Alignment with the question's directive keyword** (e.g., "Critically examine", "Discuss")

Step 3: Feedback Generation

The final report must follow this structured format:

Dear Student,

Total Marks: X.X

Introduction

[3-line comment aligned with SOP criteria]

Body

- [Comment 1]
- [Comment 2]
- [Comment 3]

Marks Awarded: X.X

Conclusion

- [Comment 1]
- [Comment 2]
- [Comment 3]

Marks Awarded: X.X

Presentation

Strengths:

- [Point 1]
- [Point 2]

Improvements:

- [Point 1]
- [Point 2]
- [Point 3]

Marks Awarded: X.X

Strengths

- You have a fair understanding of the topics.
- Your contextual understanding is really appreciable.

- You have used the Introduction – Body – Conclusion approach.
- Your handwriting is neat and legible.

All the Best! Keep improving and striving for excellence.

Tooling Requirements

Candidates can use modern tooling such as:

- **LLMs:** GPT-4o, Claude 3, Gemini, Mistral, LLaMA3
 - **OCR:** PaddleOCR, Tesseract, Vision APIs
 - **RAG Components:**
 - **Embeddings:** Cohere, OpenAI, SentenceTransformers
 - **Vector DB:** Pinecone, Qdrant, Weaviate
 - **Frameworks:** LangChain, LlamaIndex
 - **Frontend:** Streamlit, Gradio, or FastAPI
 - **Optional Enhancements:**
 - Re-rankers
 - Section-specific chunking
 - Diagram detection and markup
-

Deliverables

1. GitHub link (or ZIP) containing:
 - Code
 - Preprocessed SOP
 - Chunking and embedding strategy
 - Inference flow
 2. Example Input: 1 sample answer sheet (you may simulate)
 3. Example Output: Corresponding evaluation report
 4. README with:
 - Tech stack
 - RAG architecture diagram
 - Evaluation pipeline explanation
-

Bonus Points

- Hosted demo on Hugging Face Spaces / Streamlit Cloud
 - Hindi language handling and alignment with SOP tone
 - Vision-LLM integration (e.g., GPT-4o)
-

Timeline

Submit your task within **2 days** from assignment.
